

Symbolic Security Predicates: Hunt Program Weaknesses

Alexey Vishnyakov*, Vlada Logunova*[‡], Eli Kobrin*[†], Daniil Kuts*, Darya Parygina*[†] and Andrey Fedotov*

*Ivannikov Institute for System Programming of the RAS

[†]Lomonosov Moscow State University

[‡]Moscow Institute of Physics and Technology

Moscow, Russia

{vishnya, vlada, kobrineli, kutz, pa_darochek, fedotoff}@ispras.ru

Abstract—Dynamic symbolic execution (DSE) is a powerful method for path exploration during hybrid fuzzing and automatic bug detection. We propose security predicates to effectively detect undefined behavior and memory access violation errors. Initially, we symbolically execute program on paths that don't trigger any errors (hybrid fuzzing may explore these paths). Then we construct a symbolic security predicate to verify some error condition. Thus, we may change the program data flow to entail null pointer dereference, division by zero, out-of-bounds access, or integer overflow weaknesses. Unlike static analysis, dynamic symbolic execution does not only report errors but also generates new input data to reproduce them. Furthermore, we introduce function semantics modeling for common C/C++ standard library functions. We aim to model the control flow inside a function with a single symbolic formula. This assists bug detection, speeds up path exploration, and overcomes overconstraints in path predicate. We implement the proposed techniques in our dynamic symbolic execution tool Sydr. Thus, we utilize powerful methods from Sydr such as path predicate slicing that eliminates irrelevant constraints.

We present Juliet Dynamic to measure dynamic bug detection tools accuracy. The testing system also verifies that generated inputs trigger sanitizers. We evaluate Sydr accuracy for 11 CWEs from Juliet test suite. Sydr shows 95.59% overall accuracy. We make Sydr evaluation artifacts publicly available to facilitate results reproducibility.

Index Terms—security predicate, automatic bug detection, function semantics, symbolic execution, concolic execution, dynamic analysis, binary analysis, computer security, security development lifecycle, weakness, bug, error, sanitizer, Juliet, DSE, SMT, DAST, SDL

I. INTRODUCTION

Modern software is rapidly developing. Novel code inevitably brings new bugs and weaknesses [1]. More and more industrial companies follow security development lifecycle [2–4] to improve application quality and defend it from malicious attacks. Coverage-guided fuzzing [5–7] is continuously applied to detect program crashes during development process.

This work was supported by RFBR grant 20-07-00921 A.

Advanced hybrid fuzzing benefits from dynamic symbolic execution (DSE) [8–17]. DSE discovers complex program states that are hardly reachable via simple fuzzing.

The main focus of this paper is the promising bug-driven hybrid testing approach [18, 19]. We propose automatic error detection for undefined behaviour and memory access violation errors. In particular, we utilize dynamic symbolic execution to generate program input data that trigger integer overflow [18, 20, 21], out-of-bounds access [18], etc. These are the most widespread and dangerous program weaknesses [22]. We symbolically execute binary code because it's more portable than building from varying programming languages. The fuzzing setup may be the following. We build a target program with sanitizers [23] for fuzzer and without them for our dynamic symbolic execution tool Sydr [13]. Sydr explores new program states via branch inversion. Moreover, it generates new inputs that trigger errors. All inputs are passed to fuzzer that verifies inputs on sanitizers. Thus, fuzzer takes inputs from Sydr that increase code coverage and verifies whether any inputs crash program on sanitizers.

This paper makes the following contributions:

- We introduce symbolic function semantics for common C/C++ standard library functions. We aim to model the control flow inside a library function with a single symbolic formula. Thus, we can explore more program states from one concrete execution path. This speeds up program paths exploration and assists bug detection.
- We propose security predicates that allow us to effectively detect undefined behavior and memory access violation errors on initially valid paths. Moreover, dynamic symbolic execution produces inputs that trigger such errors. These inputs may later be verified on sanitizers [23].
- We present Juliet Dynamic [24] to measure dynamic bug detection tools accuracy on Juliet test suite [25]. We evaluate accuracy of our security predicates implementation in Sydr [13] for 11 CWEs [1] from Juliet. Sydr shows 95.59% overall accuracy.

The paper is organized as follows. Section II discusses function semantics modeling during symbolic execution. Section III presents null pointer dereference, division by zero, out-of-bounds access, and integer overflow security predicates. Section IV sheds light on implementation details. Section V contains function semantics evaluation, Juliet Dynamic [24] description, and Sydr security predicates evaluation on Juliet test suite [25]. Section VI concludes the paper.

II. FUNCTION SEMANTICS

Performance is crucial in dynamic symbolic execution. Plenty of optimizations are implemented to decrease time required for analysis. It's worth to reduce instrumentation where it is possible since the instrumented process is always slower than native execution in orders of magnitude.

There is no need to symbolically execute every internal instruction in libc functions with standard specified semantics. Moreover, cutting-edge concolic executors adapting such approach benefit from diminishing overconstraining of symbolic state [14, 16, 26, 27]. For instance, both uppercase and lowercase characters are permissible as input data for `tolower` what can be expressed in the following if-then-else (*ite*) formula for input symbol *ch*:

$$ite(ch - 'A' < 26, ch - ('A' - 'a'), ch)$$

However, relying on concrete execution trace inside this function during analysis ends up in overconstraining to letter case from concrete execution.

Instead, we propose function semantic models which can incorporate more symbolic states in a single constructed formula and speed up the execution. Similar techniques are applied in a range of notable DSE tools. In our view, there are four main directions in function semantics modeling which prevail or can be combined in a certain tool:

- 1) KLEE [28] replaces libc functions with their simpler versions precompiled to LLVM IR. This aims to curtail extra components compilation and simplify the analyzed code. This method doesn't involve symbolic modeling of entire function.
- 2) The function instrumentation is replaced with path constraints reproducing concrete execution trace [14].
- 3) The function is modeled with a single branch for representative states [16], e.g. strings are equal in `strcmp`.
- 4) Our approach generally focuses on construction of symbolically computed expression for return value [26, 27].

We handle function semantics similar to Fuzzolic analysis modes [16] except that we always perform concrete execution via dynamic binary instrumentation [29] to update concrete state. Symbolic execution starts from the first symbolic input read. We divide symbolic function models in several groups. Some of them move around already existing symbolic expressions in memory and create necessary path constraints while others just skip symbolic execution. The more complicated semantics deal with comparisons and string to integer conversions that require construction of symbolic formulas representing return values. It is worth noting that Sydr always

synchronizes concrete and symbolic states. The symbolic state is concretized if two states mismatch. Thus, we should mind that symbolic formulas should comply with the concrete state.

A. No Symbolic Computation

This section comprises functions with trivial side-effects. We handle them to increase performance and simplify the symbolic formulas.

1) *Dynamic Memory*: Heap allocation functions (`malloc`, `calloc`, `realloc`, `free`) often operate on symbolic data. We execute such functions concretely and maintain symbolic memory state consistency. Thus, we are able to radically reduce the number of overconstraint branch conditions in the path predicate. In practice, skipping symbolic execution of internal instructions for these functions provides significant increase in performance.

2) *Data Movement*: Symbolic model of copying operations (like `strcpy`, `memcpy`, `memmove`) reassigns symbolic expressions without modification. Nevertheless, we should consider the null-terminator position as it may spawn new branches.

3) *Printing Omission*: Functions which write program log to standard output are not of interest for symbolic execution process [16]. Consequently, they can be skipped in favour of performance improvement. Moreover, this might decrease the number of irrelevant constraints for the printed data.

B. String Comparison

We utilize Triton framework [9] to construct formulas for return values of functions `memchr`, `strlen`, `strcmp`, `memcmp`, `strstr`, and etc. There are two basic semantics in this group: 1) character search in string or memory area (à la `memchr`), 2) lexicographical comparison of two strings or memory areas (à la `memcmp`). The `strstr` function can be considered as `strchr` with extended character argument.

1) *à la memchr*: Formula for `memchr`-like functions should express pointer to the character in string or null-pointer when it cannot be found. Both string and character may be symbolic. We compute the character position via sum of if-then-else (*ite*) expressions that evaluate to 1 or 0. Here is the formula for `memchr(ptr, ch, count)` return value:

$$ite\left(\bigwedge_{k=0}^{count-1} ptr[k] \neq ch, 0, ptr + \sum_{i=0}^{count-1} ite\left(\bigwedge_{k=0}^i ptr[k] \neq ch, 1, 0\right)\right)$$

The first *ite*-condition checks if the character is present in string and sets the pointer to null otherwise. In the second part every summand increments the pointer until the required character is discovered. It is worth noting that `strlen(str)` is similar to `memchr(str, 0, length + 1)`.

When handling strings (e.g. `strchr`) the formula is burdened with extra conditions caused by null-terminator consideration. In case of `strstr`, character search $ptr[k] \neq ch$ is replaced with appropriate conjunctions for substring comparisons.

2) *à la memcmp*: The task of lexicographical comparison (`memcmp`, `strcmp`, `strncmp`, etc.) can be formulated as mismatch search and difference evaluation. The `memcmp(lhs, rhs, count)` formula is the following:

$$lhs[0] - rhs[0] + \sum_{i=1}^{count-1} (lhs[i] - rhs[i]) * ite\left(\bigwedge_{k=0}^{i-1} lhs[k] = rhs[k], 1, 0\right)$$

The formula is designed so that only one summand might be meaningful and others are equal to zero. We make the conjunction of all previous character pairs equality to identify the first pair of differing symbols. In string comparison it extra checks that no null byte pair was discovered before.

The mentioned above function semantics specify the returned value sign but cannot guarantee concrete and symbolic states consistency. For example, in 32-bit Linux the return values set is $\{1, 0, -1\}$ while `x86_64` uses the characters difference. For that reason, the built expression should be additionally tied to current return value to avoid concretization.

C. String to Integer Conversion

The main focus of this group is on `strtol`, `strtoul`, `strtoll`, and etc. These functions are also called inside other functions like `atoi` and `scanf("%d", &x)`. In addition, similar semantics are applied for `std::cin` reading of integer types.

The byte string parsed by `strtol` may have many configurations due to optional components and can be modeled by splitting into several parts:

[whitespaces][sign][prefix]
[valid symbols][invalid symbols]nullbyte

The set of valid symbols is defined by the *base* argument and can consist of digits and latin characters. When the *base* is null, decimal numeral system is chosen by default or it might be learned from octal and hexadecimal prefixes (0 and 0x/0X respectively). The string gets through preliminary parsing to extract parts relevant to conversion (*italic*). The same structure is preserved during new inputs generation.

$$\pm (c_n c_{n-1} \dots c_1 c_0)_b \longrightarrow x \quad (1)$$

$$a_k = ite(c_k \geq '0' \wedge c_k \leq '9' \wedge c_k < '0' + b, c_k - '0', ite(c_k \geq 'a' \wedge c_k < 'a' + b - 10, c_k - 'a' + 10, c_k - 'A' + 10)) \quad (2)$$

$$|x| = \sum_{k=0}^n a_k b^k, \quad x = ite(sign = '-', -|x|, |x|) \quad (3)$$

$$(c_k \geq '0' \wedge c_k \leq '9' \wedge c_k < '0' + b) \vee (c_k \geq 'a' \wedge c_k < 'a' + b - 10) \vee (c_k \geq 'A' \wedge c_k < 'A' + b - 10) \quad (4)$$

In order to implement the symbolic string to integer conversion (1) we constrain every character (4) to assure they lay within valid symbols according to base *b* determined

from argument value or parsing. The formula (2) depicts how we compute digit from character, i.e. classical digit symbol or alphabet letter in lower/upper case. The whole number computation result can be observed in the formula (3).

The proposed model concretizes *base* argument because symbolic base produces varying scenarios with corner cases that require lots of complex symbolic formulas. So, we don't consider symbolic *base* for the sake of model simplicity. Thus, we should also constrain the base prefix. When the *base* argument is zero, the leading zeros would cause decimal base distortion into octal. For instance, the initial input 4567 defines *base* = 10 and new input 0123 shouldn't be generated for number 123_{10} as it will be naturally identified as 123_8 . The above corner case can be solved by replacing leading zero digits with spaces. Furthermore, additional constraints are constructed to make sure the space symbols wouldn't appear between digits. However, this situation wouldn't occur if the *base* argument is a concrete non-zero value.

The obtained formula (3) may overflow, i.e. the computed value won't fit in return type (`long` for `strtol`). Therefore, we compute (3) in a bit vector of twice bigger size than return type. Afterwards, we constrain the computed value to fit in return type, e.g. $LONG_MIN \leq x \leq LONG_MAX$.

III. SECURITY PREDICATES

Dynamic symbolic execution is a powerful method to discover new program paths during hybrid fuzzing [11, 13–17]. Furthermore, we show that it can be effective to detect bugs on valid paths. We symbolically execute a program with input that doesn't lead to crash. If instruction or function operates on symbolic data, then we construct security predicates that check for undefined behavior and memory access violation. *Security predicate* for some error type (weakness) is a Boolean predicate that holds true iff the instruction (or function) triggers an error. Further, we conjunct a security predicate with sliced branch constraints from the path predicate, i.e. constraints over symbolic variables that are relevant to variables in security predicate [13]. Thus, we guarantee that program reaches the examined instruction. Finally, we call solver to resolve the constructed constraint. If it is satisfiable, then Sydr reports an error and generates new program input from obtained model. The generated input will reproduce the error in analyzed program.

In particular, Sydr can detect division by zero and null pointer dereference errors. Security predicates for them are equality of the divisor and memory dereference address respectively to zero. Let's consider how Sydr generates inputs that trigger null pointer dereference. The following example contains possible null pointer dereference in line 6:

```
1 char a[] = {'1', '2', '3', '4', '5'};
2 int main() {
3     long i;
4     scanf("%ld", &i);
5     if (i >= 5) return 1;
6     printf("%c\n", a[i]);
7 }
```

Listing 1. NULL pointer dereference (Linux x64).

We run this example under Sydr with +00000000000000000001 as standard input that is enough to generate any number that can fit in 64-bit signed long integer type. This comes from limitation of dynamic symbolic execution approach. DSE cannot increase initial concrete input size, otherwise it won't interpret missing instructions operating on the additional input bytes and symbolic model will be unsound. The symbolic input is obtained via `read` system call inside `scanf` function. Sydr creates a symbolic variable for each inputted byte. Starting from this point Sydr interprets all assembly instructions that operate on symbolic values (updates symbolic registers and memory states). The necessary concrete values are gathered from DynamoRIO [29]. Then `scanf` calls `__strtoll_internal` to convert input string to integer. Sydr wraps `strtoll` and constructs corresponding formula for string to integer conversion. Later this formula is compared with size of array `a` in line 5 and $i \geq 5$ constraint is pushed to path predicate. In line 6 the index value `i` is symbolic. So, we verify whether symbolic address $a+i$ can be zero. We conjunct the null dereference security predicate with the path predicate. The resulting predicate $i \geq 5 \wedge a+i = 0$ is satisfiable because line 5 does not consider that integer may be negative. Thus, Sydr reports null pointer dereference on `movzx eax, byte ptr [rdx + rax]` instruction, where `rax` is array `a` base address and `rdx` is index `i`. Moreover, Sydr stores new `stdin` value -0000000000006295616 that triggers zero dereference.

Null pointer dereference security predicate is able to generate input computing to actual zero address when binary is compiled without ASLR. Though it is still fast and easy solvable checker that can simply generate inputs that cause segmentation faults even when ASLR is enabled.

It is worth noting that some zero dereferences cannot be found with security predicates. They should be located with simple path discovery via symbolic execution or fuzzing. For instance, pointer in listing below isn't symbolic but can be null:

```
1 int *p = 0;
2 if (some_condition) p = malloc(4);
3 *p = 3;
```

The main advantage of bug searching with DSE is that it produces inputs to verify errors. However, true positive rate is not 100%. That's why we build two target binaries. First one is with sanitizers for fuzzing. The second one is without sanitizers for DSE since we don't want to symbolically interpret instrumentation code. In this approach fuzzer verifies whether generated inputs from security predicates crash on sanitizers.

A. Out-of-bounds Access

Out-of-bounds access is the most dangerous and widespread program error [22]. We build security predicate for each memory access at symbolic address `sym_addr` (that depends on user input) to detect such errors. Firstly, we determine memory buffer bounds $[lower_bound, upper_bound]$ and create predicate that is true when symbolic address is outside of

these bounds $sym_addr < lower_bound \vee sym_addr >= upper_bound$. Afterwards, we conjunct this predicate with sliced [13] path constraints. If solver returns a model for constructed predicate, then Sydr reports an out-of-bounds error and generates a corresponding input, i.e. original input with bytes replaced with ones from the model. It is worth noting that both bounds cannot be always determined in binary code. Thus, we try to point symbolic address outside (below or above) the single detected bound. For instance, in Listing 1 out-of-bounds is also feasible but Sydr cannot detect an upper bound for global array `a`. However, Sydr is able to heuristically retrieve the array base address from symbolic address expression $[rdx + rax]$, where `rax` is concrete array base address and `rdx` is symbolic index. Sydr assumes the concrete part `rax` to be lower bound. Thus, it generates input that triggers access below the lower bound (e.g. -1).

We maintain shadow heap and stack in order to determine symbolic address bounds. Sydr wraps all dynamic memory management functions (`malloc`, `calloc`, `realloc`, `free`, etc.) and appropriately updates the shadow heap that holds all allocated memory buffer bounds. On each call instruction Sydr pushes return address location (stack pointer value) to shadow stack. And on all call and return instructions Sydr pops elements from shadow stack according to current stack pointer value.

When instruction accesses memory at symbolic address Sydr detects corresponding buffer bounds the following way. If current concrete address value is in shadow heap, then both bounds are retrieved from it. When address points to stack, the closest return address location (from shadow stack) above concrete address is the upper bound. The lower bound is computed heuristically from the concrete part of symbolic address formula. The main idea behind this approach is to sum up the concrete parts of the formula. However, it considers some corner cases. For instance, we should distinguish $a[i - 0x20]$ from array on stack $[ebp - 0x20]$. We use the same heuristics to detect global array lower bound.

Moreover, Sydr wraps memory copy functions (`memcpy`, `memmove`, `memset`, `strncpy`, etc.) to detect buffer overflows. If memory copy size function argument is symbolic, Sydr tries to make it exceed the upper bound.

Before solving the security predicate Sydr conjuncts it with strong precondition to make error most likely cause a crash, i.e. overwrite return address or dereference negative address. If such predicate is unsatisfiable, Sydr falls back to solving the original security predicate. Furthermore, when additionally both address and value are symbolic, Sydr reports that write-what-where condition is possible for this out-of-bounds access.

B. Integer Overflow

Integer overflow is one of the most common program errors [22]. However, it occurs quite often in binary code. Thus, Sydr would operate too long if we check all the cases when integer overflow may happen. Moreover, there are some situations like hash functions where integer overflow is normal.

Algorithm 1 Signedness detection via backward slicing.

Input: *sink* – sink AST node containing overflowed value, *call_stack* – call stack on error sink, Π – path predicate (path constraints prior to the error sink).

Output: Return True if type is signed, False if unsigned, None if cannot be detected.

```
vars ← used_variables(sink)           ▷ slicing variables
▷ Iterate backward over path constraints.
for all c ∈ reversed( $\Pi$ ) do
  ▷ Check whether constraint uses sink node vars.
  if vars ∩ used_variables(c) ≠ ∅ then
    ▷ Get instruction corresponding to constraint c.
    inst ← instruction(c)
    ▷ Get call site of function containing inst.
    s ← callsite(function(inst))
    if inst is branch and s ∈ call_stack then
      if inst is signed then           ▷ js/jns/jg/jge/jl/jle
        return True
      if inst is unsigned then         ▷ ja/jae/jb/jbe
        return False
    ▷ Signedness may be ambiguous (e.g. jz).
```

```
return None
```

So, we highlight only critical parts and verify security predicates for them. Unlike other security predicates, we separate an error sink from its source. Source is an instruction where integer overflow may happen. And sink is a place in code where preceding flaw may lead to critical error.

We solve integer overflow security predicates in error sinks that use potentially overflowed value, i.e. branches (changing control flow depending on overflowed value), memory access addresses, and function arguments. This is especially critical for such functions as `malloc`, `memcpy`, and others [20, 21]. We wrap some common standard library functions and consider all their symbolic arguments to be potential error sinks. For other function calls we check the first three arguments according to standard calling convention.

First of all, we check whether instruction is arithmetic and one of its operands is symbolic. If so, we build security predicates for unsigned (CF) and signed (OF) integer overflow errors. For most arithmetic instructions security predicate is true when the corresponding flag is set to 1.

Then we figure out whether error source instruction result is involved in computation of sink. Firstly, we check whether at least one children of sink AST matches the source AST (overflowed arithmetic instruction result). If so, sink contains potentially overflowed value in its computations.

Afterwards, we identify the signedness of arithmetic operation. So, we could select one of signed or unsigned integer overflow security predicates. It is important because the lack of signedness knowledge leads to many false positives. We propose Algorithm 1 to detect signedness via backward slicing [30]. We iterate backwards starting from the last branch in path predicate (at the moment of error sink). We locate

the first branch that we can learn signedness from [20]. For example, `JL` branch instruction tells us that value is signed. Moreover, the branch must use at least one symbolic variable from sink AST and its call site have to be in current call stack, i.e. function containing this branch should be in current call stack.

Furthermore, we can detect signedness if symbolic value is obtained with `strtol` functions. For instance, `strtol` is used for signed value and `strtoul` for unsigned value.

Finally, when we get information about the signedness (or do not), we check whether security predicates are satisfiable. If we know the signedness, we check only one specific security predicate for this signedness. We slice all relevant constraints from the path predicate at sink point and conjunct them with the security predicate. We report signed or unsigned integer overflow and generate corresponding input when the resulting predicate is satisfiable.

If we couldn't retrieve the signedness, then we query whether the integer overflow of both kinds may happen at the same time. In other words, may some single input cause unsigned and signed integer overflow errors simultaneously. So, we conjunct security predicates of both signednesses and sliced path predicate. If the resulting predicate is true, we report both signed and unsigned overflows and generate a single corresponding input. Otherwise, we check security predicates separately. If both security predicates are separately true, we print two warnings and save two inputs for each of them. We do not report errors when single signed or unsigned integer overflow is possible without knowing the signedness since it may lead to false positives.

As we said before, it is crucial when potential sink is argument of functions like `malloc`, `calloc`, `memcpy`, etc. Thereby, we do not only check for ordinary integer overflow but we reveal whether even more dangerous consequences might appear. For allocation functions we examine whether the size can be overflowed such that it would be a non-zero value smaller than it was during the concrete execution. In particular, it may further lead to buffer overflow. For copying functions we investigate whether the size can be overflowed such that it would be greater than its value during the concrete execution. For such cases we build additional strong preconditions that are conjuncted with security predicate. If conjunction with precondition is unsatisfiable, we fall back to solving the original security predicate.

Let us consider the example of program with potential integer overflow that may lead to buffer overflow on 32-bit architecture:

```
1 int main() {
2   int size;
3   fscanf(stdin, "%d", &size);
4   if (size <= 0) return 1;
5   size_t i;
6   int *p = malloc(size * sizeof(int));
7   if (p == NULL) return 1;
8   for (i = 0; i < (size_t)size; i++) {
9     p[i] = 0;
10  }
```

```

11     printf("%d\n", p[0]);
12     free(p);
13 }

```

Listing 2. Integer overflow in `malloc` size (32-bit).

We run this program under Sydr with standard input +00000000002 that is enough to generate 32-bit positive and negative numbers. The program gets the size of the array to allocate via `strtol` function called from `scanf`. String to integer conversion function semantics are applied to obtain symbolic formula for `size`. The array `size` is multiplied by `sizeof(int)` which is a potential error source. We construct signed integer overflow security predicate for the multiplication (signedness is identified from `strtol` call). The potential error sink is in `malloc` function argument that uses the multiplied value. Moreover, we conjunct the security predicate with strong precondition that $size * sizeof(int)$ should be lower than original $2 * 4 = 8$. This precondition helps to bypass check in line 7, i.e. we shouldn't allocate more than total memory available on the machine. Then we conjunct security predicate, precondition, and sliced path predicate. The obtained predicate $(long\ long)size * sizeof(int) \neq (int)size * sizeof(int) \wedge (unsigned)size * sizeof(int) < 8 \wedge size > 0$ is passed to solver. The solver reports it's satisfiable and returns a corresponding model. Sydr generates new standard input value +01073741825 that triggers integer overflow in line 6. And `malloc` allocates only $(unsigned)1073741825 * sizeof(int) = 4$ bytes because of wraparound. Thus, in line 9 we get buffer overflow since cycle iterates over 1073741825 array elements when its real size is 1.

1) *Error Source Pitfalls:* We build security predicates for SHL and SAL instructions differently because CF flag in these instructions contains the value of the last bit shifted out of the destination operand and OF flag is affected only for 1-bit shifts. So, CF and OF flags do not represent the fact that overflow occurred. Therefore, unsigned integer overflow security predicate is true when at least one of significant bits was shifted from destination operand. Signed integer overflow occurs when both zeros and ones were shifted.

On 32-bit architecture `int64_t` addition and subtraction are implemented as a sequence of ADD/ADC or SUB/SBB instructions. Therefore, we build the security predicate only for the last ADC/SBB instruction in sequence.

Moreover, some compiler optimizations replace subtraction of `int64_t` and unsigned `int` types with addition, e.g. replace `sub eax, 1` with `add eax, 0xffffffff` [20]. When we see an addition with negative constant, we construct a security predicate just like for subtraction with positive value.

2) *Error Sink Pitfalls:* We intend to correctly handle integer promotion. If arithmetic operations on types smaller than `int` occur, their operands get extended to bigger types. For instance, two `char` operands would be extended to `int` and addition would compute 32-bit values instead of 8-bit. Thus, overflow can never happen in the moment of addition because 32-bit can fit any result of 8-bit arithmetics. The actual error will be on the step when we truncate the result to the smaller

type back (`char a, b, c; c = a + b;`). In terms of SMT [31], the addition result is extracted to fit in 8 bits. Therefore, we search for all extraction nodes in the sink AST that extract from the source AST. Then we reconstruct new security predicates for smaller operands with sizes equal to extraction node size.

IV. IMPLEMENTATION

We implement the proposed function semantics and security predicates in our dynamic symbolic execution tool Sydr [13]. Sydr utilizes Triton [9] to construct symbolic formulas, maintain symbolic registers and memory states, and build path predicate. The obtained formulas are solved by Z3 [32].

Sydr is separated in two processes: concrete and symbolic executors. Concrete executor runs target program under dynamic binary instrumentation framework DynamoRIO [29] and sends events to symbolic executor via shared memory. These events contain all necessary instructions, registers, and memory values to perform concolic execution. We handle function semantics the similar way. DynamoRIO wraps library functions. So, we send event containing function name, its arguments, and return value for each function mentioned in Section II. Then we handle this event on symbolic executor side and apply appropriate symbolic semantics. Symbolic executor updates the shadow heap every time it receives an event for heap library functions. The shadow stack is renewed when call/ret instructions are received.

We use xxHash [33] fast hash algorithm in order to skip already discovered program errors. So, we update bitmap with all satisfiable security predicates and don't solve them again. The bit index in this bitmap corresponds to hash over error type, source, and destination addresses. Moreover, we remember path constraint index when security predicate is unsatisfiable. Thus, we can skip solving predicates with greater indexes because they are going to be unsatisfiable too.

The implementation supports x86 (32 and 64 bit) architecture. However, the proposed methods are not limited to x86 and may later be implemented for other architectures.

V. EVALUATION

A. Function Semantics

We evaluate function semantics performance and efficiency on 64-bit Linux programs [34]. We utilize AMD EPYC 7702 (128 cores) server with 256G RAM for Sydr benchmarking. Table I presents the benchmarking results for default symbolic execution and one with function semantics applied. Sydr symbolically interprets one execution trace for each program.

First of all, we measure path predicate construction time and number of discovered symbolic branches for each application. As we can see, both measured metrics decrease when modeling function semantics. Thus, we win more time for exploring new paths via branch inversion. The fewer symbolic branches we have the less overconstraint we achieve in formulas. Moreover, we do not invert internal branches in wrapped library functions and model them in a single (or few) formula. So, we are able to explore more program logic states

TABLE I
FUNCTION SEMANTICS BENCHMARKING

Application	Path predicate				2-hour benchmark							
	Default		Function Semantics		Accuracy	Default		Time	Accuracy	Function Semantics		
	Branches	Time	Branches	Time		SAT	Queries			SAT	Queries	Time
bzip2recover	5131	6s	5131	6s	100%	2101	5131	47m35s	100%	2101	5131	45m38s
cjpeg	8008	19s	6992	18s	100%	50	2656	120m	100%	50	3750	120m
faad	470585	21m	466697	15m52s	97.11%	1974	3072	120m	98.91%	1560	2414	120m
foo2lava	910737	21m9s	905592	18m20s	87.1%	31	5998	120m	99.02%	205	6668	120m
hdp	66070	43s	29265	20s	76.69%	1171	4122	120m	72.22%	5893	12172	120m
jasper	837643	14m47s	771806	10m37s	99.62%	8457	22538	120m	96.61%	9528	24472	120m
libxml2	53400	40s	8873	12s	51.27%	1063	18485	120m	82.44%	1247	8970	5m53s
minigzip	8977	1m4s	8977	1m3s	51.47%	7569	8977	16m16s	51.47%	7569	8977	16m16s
muraster	7102	5s	4453	4s	99.94%	3304	6041	120m	100%	360	470	120m
pk2bm	3665	2s	658	1s	99.45%	183	3664	15m55s	100%	189	657	4m55s
pnmhistsmap_pgm	967187	9m21s	967155	9m2s	99.99%	19351	28932	120m	100%	19964	29369	120m
pnmhistsmap_ppm	7864	12s	7822	11s	99.07%	107	7990	27m26s	99.12%	114	7948	25m31s
readelf	62713	41s	13649	10s	87.38%	1022	9541	120m	85.82%	2363	6541	120m
yices-smt2	19352	17s	10340	11s	73.79%	4258	16222	120m	70.27%	5534	11753	11m5s
yodl	8329	9s	5340	5s	36.25%	1153	9403	51m3s	98.26%	1150	6414	1m50s

from one execution trace. For instance, we can make strings equal during one run in `memcmp` function.

Secondly, we run Sydr to invert branches in one solving thread. We limit Sydr execution time with 2 hours. Each query solving time is limited to 10 seconds. In our experiment Sydr inverts branches in direct order (from first to last in path predicate). Finally, we measure the accuracy [13] of branch inversion because not every satisfiable (SAT) solver query actually inverts the target branch. *Accuracy* is the percent of satisfiable solver queries that successfully discover the intended path, i.e. they have the same execution branch trace as original except the last branch, that should be in inverted direction. For most applications Sydr shows the best results with function semantics applied. Sydr either inverts all branches on execution trace faster, or discovers more (accurate) paths for the same 2-hour limit. The number of discovered paths for `faad` and `muraster` decrease because they contain lots of symbolic string comparisons. We model each string comparison function with a single data flow formula. Thus, we skip branches inside these functions. These branches are useless since they compare single bytes. Instead, we perform full string comparison. We manually verified that we actually invert complex branches that use string comparison return values. For example, `if (!strcmp(s, "abc"))`. However, some solver queries become more complex because they contain sliced [13] (conjoined) string comparison branch constraints. This fact negatively affects the solver performance.

B. Security Predicates on Juliet

Juliet Test Suite for C/C++ [25] is a collection of test cases for 118 different CWEs. We adopted Juliet build system to make it suitable for dynamic analysis. We build only those test cases which read user input, i.e. it can be modified to cause a program error. Each Juliet source file contains two test cases: positive and negative. Positive test case has an error while negative case handles potential flaw via additional checks. These cases are implemented in `_bad` and `_good` functions that are called directly from `main`. We build these tests in separate binaries: one with potential flaw and other without

errors. Thus, each binary runs a single (positive or negative) test case. Furthermore, for each binary we build its version with sanitizers [23]. We run Sydr security predicates on binary without sanitizers. Then we verify that Sydr generates input that actually triggers an error on sanitizers. We propose Juliet C/C++ Dynamic Test Suite [24] that allows to measure true positive and true negative rates for a dynamic error detection tool.

It is worth noting that Juliet test cases read symbolic input with `fscanf` or `fgets+atoi`. Both use `strtol` functions for string to integer conversion, which makes function semantics essential for error detection in Juliet tests.

The evaluation results are presented in Table II. We measure Sydr on a subset of Juliet CWEs. Out-of-bounds security predicate cover CWE 121, 122, 124, 126, 127, 194, and 195. Integer overflow checker detects errors for CWE 190, 191, and 680. CWE369 is handled by division by zero predicate. We do not build some CWEs (like CWE476: NULL Pointer Dereference) because they have no user input. Every covered CWE is built for both 32-bit and 64-bit targets. The only exception is CWE680 (Integer Overflow to Buffer Overflow) that is built only for 32-bit. For instance, in Listing 2 input variable `size` is a 32-bit integer. Potential integer overflow happens in `malloc` function argument. Inputted `size` is multiplied by `sizeof(int)` that has `size_t` type. On 64-bit architecture `size_t` is equivalent for unsigned long long and 32-bit value, sign-extended to 64-bit, cannot cause integer overflow.

Juliet Dynamic testing system evaluates the number of true positive (*TP*) and true negative (*TN*) cases. Afterwards, it computes true positive rate $TPR = \frac{TP}{P}$, true negative rate $TNR = \frac{TN}{N}$, and accuracy $ACC = \frac{TP+TN}{P+N}$ (Juliet has equal number of positive and negative cases $P = N$). We collect these results for two categories of alarmed errors (Table II):

- Textual errors – Sydr reports that test contains an error.
- Sanitizers verification – generated inputs for errors from previous category trigger sanitizers.

We describe evaluation process below.

TABLE II
JULIET TESTING RESULTS

CWE	P=N	Textual errors			Sanitizers verification		
		TPR	TNR	ACC	TPR	TNR	ACC
CWE121: Stack Based Buffer Overflow	188	100%	100%	100%	100%	100%	100%
CWE122: Heap Based Buffer Overflow	376	100%	100%	100%	100%	100%	100%
CWE124: Buffer Underwrite	188	100%	100%	100%	100%	100%	100%
CWE126: Buffer Overread	188	100%	100%	100%	100%	100%	100%
CWE127: Buffer Underread	188	100%	100%	100%	100%	100%	100%
CWE190: Integer Overflow	2580	99.92%	90.89%	95.41%	98.10%	90.89%	94.50%
CWE191: Integer Underflow	1922	99.90%	91%	95.45%	97.45%	91%	94.22%
CWE194: Unexpected Sign Extension	752	100%	100%	100%	100%	100%	100%
CWE195: Signed to Unsigned Conversation Error	752	99.87%	100%	99.93%	99.87%	100%	99.93%
CWE369: Divide by Zero	564	66.67%	100%	83.33%	66.67%	100%	83.33%
CWE680: Integer Overflow to Buffer Overflow	188	100%	100%	100%	100%	100%	100%
TOTAL	7886	97.55%	94.83%	96.19%	96.36%	94.83%	95.59%

Firstly, we collect paths for test case binaries and select appropriate inputs. We crafted inputs for `int64_t`, `int`, `short`, and `char` input data types. For instance, we pass `+00000000002` to standard input for `int` type. The plus sign (`+`) allows Sydr to change the number sign. Extra zeros are required to let Sydr pick up big numbers because DSE cannot increase the number of symbolic input bytes. However, we need additional first extra zero right after the plus sign since `scanf` function makes first zero concrete:

```
if (buf[i] == '0') buf[i] = '0';
```

Secondly, we run Sydr on each test case with corresponding inputs. If Sydr reports an error for positive test, then the result is classified as true positive. If Sydr alarms on negative test, then it is a false positive. We suppose that the result is true negative when Sydr reports no errors for negative test. The false negative result happens when Sydr does not alarm an error on a positive test.

Finally, we verify generated inputs on previous step with sanitizers. Sydr generates one or multiple inputs that should reproduce a true positive error. For instance, there may be multiple errors for nested function sinks. We run all generated inputs for a single test on sanitizers. If at least one input leads to sanitizers error detection, it stays as true positive. Otherwise, we change it to false negative after sanitizers verification. False positive and true negative results stay the same.

Table II contains measured *TPR*, *TNR*, *ACC* for textually alarmed errors and their verification on sanitizers. We conclude that Sydr has 95.59% overall accuracy for 11 CWEs from Juliet. Sydr completely covers test suites for CWE 121, 122, 124, 126, 127, 194, and 680. Security predicates miss some division by zero (CWE369) errors because Sydr is based on Triton symbolic engine [9] that doesn't support floating point. As you can see, integer overflow and underflow (CWE190/191) have some false positive and false negative errors. On 32-bit architecture `int64_t` type produces large number arithmetics that our security predicates support for addition and subtraction. However, we haven't implemented multiplication yet. Other problem is caused by subtractions of `int64_t` on 32-bit architecture. Some compiler optimizations replace subtraction of `int64_t` types with addition,

e.g. replace `sub eax, 1` with `add eax, 0xffffffff`. We support this for regular arithmetics but not for large arithmetics. Moreover, there are some wrong results in test cases containing squaring of `short` type that need further investigation.

Sydr evaluation artifacts are available in Juliet Dynamic repository [24]. Thus, one can run the provided script to reproduce the results.

VI. CONCLUSION

We propose security predicates that enable bug detection in our dynamic symbolic execution tool Sydr [13]. In particular, Sydr is able to report null pointer dereference, division by zero, out-of-bounds access, and integer overflow errors. Unlike static analysis tools, Sydr also generates new input data that reproduce detected errors. We present function semantics modeling for common C/C++ standard library functions that assists bug detection, increases path exploration speed, and reduces the number of overconstrained formulas. Last but not least, we introduce Juliet Dynamic testing system [24] that allows to measure dynamic bug detection tools accuracy on Juliet test suite [25]. Sydr achieves 95.59% overall accuracy for 11 CWEs from Juliet.

REFERENCES

- [1] *CWE - common weakness enumeration*. [Online]. Available: <https://cwe.mitre.org/>.
- [2] M. Howard and S. Lipner, *The security development lifecycle*. Microsoft Press Redmond, 2006, vol. 8. [Online]. Available: <http://msdn.microsoft.com/en-us/library/ms995349.aspx>.
- [3] *ISO/IEC 15408-3:2008: Information technology – security techniques – evaluation criteria for it security – part 3: Security assurance components*, ISO Geneva, Switzerland, 2008. [Online]. Available: <https://www.iso.org/standard/46413.html>.
- [4] *GOST R 56939-2016: Information protection. secure software development. general requirements*, National Standard of Russian Federation, 2016. [Online]. Available: <http://protect.gost.ru/document.aspx?control=7&id=203548>.

- [5] K. Serebryany, "Continuous fuzzing with libFuzzer and AddressSanitizer," in *2016 IEEE Cybersecurity Development (SecDev)*, IEEE, 2016, p. 157. DOI: 10.1109/SecDev.2016.043.
- [6] S. Sargsyan, J. Hakobyan, M. Mehrabyan, M. Mishechkin, V. Akozin, and S. Kurmangaleev, "ISP-Fuzzer: Extendable fuzzing framework," in *2019 Ivannikov Memorial Workshop (IVMEM)*, IEEE, 2019, pp. 68–71. DOI: 10.1109/IVMEM.2019.00017.
- [7] A. Fioraldi, D. Maier, H. Eißfeldt, and M. Heuse, "AFL++: Combining incremental steps of fuzzing research," in *14th USENIX Workshop on Offensive Technologies (WOOT 20)*, 2020. [Online]. Available: <https://www.usenix.org/system/files/woot20-paper-fioraldi.pdf>.
- [8] J. C. King, "Symbolic execution and program testing," *Communications of the ACM*, vol. 19, no. 7, pp. 385–394, 1976. DOI: 10.1145/360248.360252.
- [9] F. Saudel and J. Salwan, "Triton: A dynamic symbolic execution framework," in *Symposium sur la sécurité des technologies de l'information et des communications*, ser. SSTIC, 2015, pp. 31–54. [Online]. Available: https://triton.quarkslab.com/files/sstic2015_slide_en_saudel_salwan.pdf.
- [10] N. Stephens, J. Grosen, C. Salls, A. Dutcher, R. Wang, J. Corbetta, Y. Shoshitaishvili, C. Kruegel, and G. Vigna, "Driller: Augmenting fuzzing through selective symbolic execution," in *NDSS*, vol. 16, 2016, pp. 1–16.
- [11] I. Yun, S. Lee, M. Xu, Y. Jang, and T. Kim, "QSYM: A practical concolic execution engine tailored for hybrid fuzzing," in *27th USENIX Security Symposium*, 2018, pp. 745–761. [Online]. Available: <https://www.usenix.org/system/files/conference/usenixsecurity18/sec18-yun.pdf>.
- [12] R. Baldoni, E. Coppa, D. C. D'Elia, C. Demetrescu, and I. Finocchi, "A survey of symbolic execution techniques," *ACM Computing Surveys*, vol. 51, no. 3, 2018. DOI: 10.1145/3182657.
- [13] A. Vishnyakov, A. Fedotov, D. Kuts, A. Novikov, D. Parygina, E. Kobrin, V. Logunova, P. Belecky, and S. Kurmangaleev, "Sydr: Cutting edge dynamic symbolic execution," in *2020 Ivannikov ISPRAS Open Conference (ISPRAS)*, IEEE, 2020, pp. 46–54. DOI: 10.1109/ISPRAS51486.2020.00014.
- [14] S. Poeplau and A. Francillon, "Symbolic execution with SymCC: Don't interpret, compile!" In *29th USENIX Security Symposium (USENIX Security 20)*, 2020, pp. 181–198. [Online]. Available: <https://www.usenix.org/system/files/sec20-poeplau.pdf>.
- [15] —, "SymQEMU: Compilation-based symbolic execution for binaries," in *Proceedings of the 2021 Network and Distributed System Security Symposium*, 2021. DOI: 10.14722/ndss.2021.23118.
- [16] L. Borzacchiello, E. Coppa, and C. Demetrescu, "FUZZOLIC: Mixing fuzzing and concolic execution," *Computers & Security*, vol. 108, p. 102368, 2021. DOI: 10.1016/j.cose.2021.102368.
- [17] D. Kuts, "Towards symbolic pointers reasoning in dynamic symbolic execution," in *2021 Ivannikov Memorial Workshop (IVMEM)*, IEEE, 2021. [Online]. Available: <https://arxiv.org/abs/2109.03698>.
- [18] Y. Chen, P. Li, J. Xu, S. Guo, R. Zhou, Y. Zhang, T. Wei, and L. Lu, "SAVIOR: Towards bug-driven hybrid testing," in *2020 IEEE Symposium on Security and Privacy (SP)*, 2020, pp. 1580–1596. DOI: 10.1109/SP40000.2020.00002.
- [19] S. Österlund, K. Razavi, H. Bos, and C. Giuffrida, "ParmeSan: Sanitizer-guided greybox fuzzing," in *29th USENIX Security Symposium (USENIX Security 20)*, 2020, pp. 2289–2306. [Online]. Available: <https://www.usenix.org/system/files/sec20-osterlund.pdf>.
- [20] T. Wang, T. Wei, Z. Lin, and W. Zou, "IntScope: Automatically detecting integer overflow vulnerability in x86 binary using symbolic execution," in *NDSS*, 2009.
- [21] R. Demidov, A. Pechenkin, and P. Zegzhda, "Integer overflow vulnerabilities detection in software binary code," in *Proceedings of the 10th International Conference on Security of Information and Networks*, ser. SIN '17, ACM, 2017, pp. 101–106. DOI: 10.1145/3136825.3136872.
- [22] *2021 CWE top 25 most dangerous software weaknesses*. [Online]. Available: https://cwe.mitre.org/top25/archive/2021/2021_cwe_top25.html.
- [23] D. Song, J. Lettner, P. Rajasekaran, Y. Na, S. Volckaert, P. Larsen, and M. Franz, "SoK: Sanitizing for security," in *2019 IEEE Symposium on Security and Privacy (SP)*, 2019, pp. 1275–1295. DOI: 10.1109/SP.2019.00010.
- [24] *Juliet C/C++ dynamic test suite*. [Online]. Available: <https://github.com/ispras/juliet-dynamic>.
- [25] *Software assurance reference dataset - Juliet test suite*. [Online]. Available: <https://samate.nist.gov/SRD/testsuite.php>.
- [26] V. Chipounov, V. Kuznetsov, and G. Candea, "The S2E platform: Design, implementation, and applications," *ACM Transactions on Computer Systems (TOCS)*, vol. 30, no. 1, pp. 1–49, 2012. DOI: 10.1145/2110356.2110358.
- [27] Y. Shoshitaishvili, R. Wang, C. Salls, N. Stephens, M. Polino, A. Dutcher, J. Grosen, S. Feng, C. Hauser, C. Kruegel, and G. Vigna, "SOK: (state of) the art of war: Offensive techniques in binary analysis," in *2016 IEEE Symposium on Security and Privacy (SP)*, 2016, pp. 138–157. DOI: 10.1109/SP.2016.17.
- [28] C. Cadar, D. Dunbar, and D. R. Engler, "KLEE: Unassisted and automatic generation of high-coverage tests for complex systems programs," in *OSDI*, vol. 8, 2008, pp. 209–224. [Online]. Available: https://static.usenix.org/events/osdi08/tech/full_papers/cadar/cadar.pdf.
- [29] D. Bruening, "Efficient, transparent, and comprehensive runtime code manipulation," Ph.D. dissertation, Massachusetts Institute of Technology, Department of Electrical Engineering and Computer Science, 2004. [On-

- line]. Available: <https://www.burningcutlery.com/derek/docs/phd.pdf>.
- [30] M. Weiser, "Program slicing," *IEEE Transactions on Software Engineering*, vol. SE-10, no. 4, pp. 352–357, 1984. DOI: 10.1109/TSE.1984.5010248.
- [31] C. Barrett, P. Fontaine, and C. Tinelli, *The SMT-LIB Standard: Version 2.6*, 2017. [Online]. Available: www.SMT-LIB.org.
- [32] L. de Moura and N. Bjørner, "Z3: An efficient SMT solver," in *Tools and Algorithms for the Construction and Analysis of Systems*, Springer Berlin Heidelberg, 2008, pp. 337–340. DOI: 10.1007/978-3-540-78800-3_24.
- [33] *xxHash - extremely fast non-cryptographic hash algorithm*. [Online]. Available: <http://www.xxhash.com/>.
- [34] *Sydr benchmark*. [Online]. Available: <https://github.com/ispras/sydr-benchmark>.